



A BUYER'S GUIDE TO AI SEARCH ACCOUNTABILITY

The GEO Accountability Checklist

How to verify your agency's AI search claims
— instead of trusting **vanity metrics**

Hassam Shabbir

Founder & CEO, SitesBrand.com

SITESBRAND — PSYCHOLOGY-DRIVEN AI SEARCH & GROWTH AGENCY



Copyright & Method Note

The GEO Accountability Checklist: How to Verify Your Agency's AI Search Claims (Instead of Trusting Vanity Metrics)

First Edition — July 2026

Copyright © 2026 Hassam Shabbir / SitesBrand. All rights reserved. No part of this publication may be reproduced or distributed without prior written permission, except for brief quotations with attribution.

How this book handles evidence. Every external claim in this book carries a numbered citation — [1], [2], and so on — resolving to a full URL, publisher, and access date in the Sources section at the end. Where a statement is our professional reasoning rather than a cited fact, it is labeled as such in the text. Nothing in this book is supported by an invented statistic, and no source is cited that we did not open and read on the access date shown.

About the case-study data. The Trackora, Finovo, and Healthify figures used in Chapter 3 are reproduced exactly as published on SitesBrand's public case-studies page [9], which notes that client data is anonymized for confidentiality. They are presented here as an example of *outcome-format reporting* — not as a promise of results for any other business.

Published by: SitesBrand — sitesbrand.com

Free audit: sitesbrand.com/free-audit

Who This Book Is For

This is a short book with one job: to help you, the **paying client**, tell the difference between real AI-search work and confident-sounding noise.

You should read it if:

- You currently pay an agency or freelancer for SEO, "AI SEO," AEO, or GEO work and receive reports you can't independently check.
- You're evaluating proposals that promise visibility in ChatGPT, Google AI Overviews, Perplexity, or Gemini.
- You've been shown a screenshot of an AI mentioning a brand and asked to treat it as proof.

What this book is not. It is not a how-to-do-GEO manual. That space is well served — including by frameworks from agencies and platforms such as Insightland's CAT framework [12] and OptimizeGEO's step-by-step guides [13], and by our own 58-page *Ultimate 2026 Guide to Modern Search Optimization* [14], which covers SEO, AEO, GEO, AIO, and SXO as one system. This book deliberately covers the opposite side of the table: **how to verify someone else's GEO work when you're the one paying for it.**

Why we wrote it. GEO measurement is genuinely immature as a discipline — the evidence for that is laid out, with sources, in Chapter 1. Immature measurement plus client urgency is historically the exact recipe that produced the worst of the SEO industry's snake-oil era. Rand Fishkin of SparkToro put it bluntly in his January 2026 research on AI answer inconsistency: "AI visibility 'experts' could easily weaponize this knowledge, just as sketchy SEO salespeople did in decades past." [5]

We run an AI search optimization practice ourselves [11]. We are writing the verification manual we believe every client of every agency — including ours — should be holding.

— **Hassam Shabbir**, Founder & CEO, SitesBrand

Table of Contents

Copyright & Method Note	2
Who This Book Is For	3
Chapter 1: Why GEO Reporting Is Easy to Fake Right Now	5
Chapter 2: The Vanity-Metric Checklist — Report Lines That Sound Good but Prove Nothing	8
Chapter 3: What a Real GEO Report Contains	11
Chapter 4: Five Questions to Ask Any Agency (Before Signing — or Mid-Contract)	16
The One-Page GEO Accountability Checklist	19
Sources	20

Chapter 1: Why GEO Reporting Is Easy to Fake Right Now

Key Takeaways

- There is **no Google Search Console equivalent for GEO** — no authoritative, engine-provided data source a client can use to double-check an agency's claims across AI platforms [7].
 - AI answers are **non-deterministic**: the same prompt produces materially different brand lists run to run, which makes single-run "AI rankings" close to meaningless [5].
 - Even the paid tracking tools **disagree with each other** on identical prompts, and practitioners quoted in trade press openly question their reliability [6].
 - Google's own reporting for generative AI features, launched June 2026, shows **impressions only — no click data** — and is still rolling out to a subset of sites [3][4].
 - None of this means GEO is fake. It means GEO **verification** has to be designed deliberately — which is what the rest of this book is for.
-

1.1 The Ground-Truth Gap

In classic SEO, an agency's claims are checkable in minutes. Google Search Console shows your clicks, impressions, and queries from Google's own logs; GA4 shows what those visitors did; any rank tracker can confirm a position claim. The data may be imperfect, but there is a shared, engine-provided source of truth.

GEO has no such thing. As the marketing-analytics company Omnia summarized in May 2026: GEO produces **no equivalent of Google Search Console** — no authoritative, real-time data source that records when your brand appears inside AI answers [7]. When an agency tells you "your brand now appears in ChatGPT answers for your category," there is no official console where you can look that up.

Google took a first step in June 2026, announcing dedicated **Search Generative AI performance reports in Search Console** — impressions data for AI Overviews and AI Mode, by page, country, device, and date [3]. Two limitations matter for accountability:

1. **No click data is included** in the current version of the report [4].
2. Google is "rolling these reports out to a subset of websites" first — your property may not have the report yet [3].

And that covers *Google's* AI features only. For ChatGPT, Perplexity, Claude, and Gemini, there is no engine-provided reporting at all as of this writing. *(That last sentence is our own assessment as of July 4, 2026, based on the absence of any such console from those vendors — not a cited fact.)*

1.2 The Non-Determinism Problem

Here is the single most important technical fact a GEO client needs to know: **AI engines do not give the same answer twice.**

In January 2026, SparkToro published the largest public study of this to date: 600 volunteers ran 2,961 prompts across ChatGPT, Claude, and Google's AI, with each of 12 B2C prompts run 60–100 times per engine [5]. The findings:

- There is **less than a 1-in-100 chance** that ChatGPT or Google's AI will produce the identical list of brands in two separate runs of the same prompt [5].
- The probability of getting the same list *in the same order* was roughly **1 in 1,000** [5].
- Even the number of recommendations varied run to run — from 2–3 items to 10+ [5].

The study's conclusion: tracking "ranking positions" in AI answers is useless; the only reasonable metric is **visibility percentage across dozens-to-hundreds of repeated runs** — how often a brand appears, not where it appeared once [5].

Now consider what this means for the report on your desk. A screenshot showing your brand as "the #2 recommendation in ChatGPT" is a photograph of a dice roll. Run the same prompt again and the list will almost certainly differ. An agency can also do the reverse: keep re-running a prompt until your brand happens to appear, screenshot that run, and present it as the result.

◆ **Our reasoning, stated plainly:** we are not claiming any specific agency does this. We are claiming the measurement environment currently *permits* it, cheaply and undetectably — unless the client knows what to demand instead. Chapter 3 covers exactly that.

1.3 Even the Professional Tools Disagree

A reasonable objection: "My agency uses a paid AI-visibility platform, so the data is solid." The trade press coverage of these tools suggests more caution.

In May 2026, Digiday reported growing marketer skepticism about AI visibility tools, citing inconsistent results, hallucination and misattribution risks, a lack of standardized benchmarks, and "point in time" measurement [6]. Paul Dyer, CEO of the AI-native services agency /prompt, put it bluntly: **"If you use three different tools and give them the same prompts, you get three different answers."** [6] One agency practitioner quoted in the same piece called the tools "just a benchmarker in my mind" [6].

These tools are not cheap, either. Digiday lists Profound at \$99/month (starter, ChatGPT tracking, 50 prompts) to \$399/month (three answer engines, 100 prompts), and Ahrefs Brand Radar at \$129–\$449/month [6]. Paying for a dashboard does not settle the question of whether the numbers on it mean anything — that depends entirely on *methodology*: how many runs, which prompts, which engines, what baseline. That's a question you can ask your agency directly (Chapter 4, Question 2).

1.4 Why This Combination Is Dangerous for Buyers

Put the three facts together:

Fact	Consequence for you
No engine-provided ground truth [7]	You cannot independently audit most claims
Answers vary run to run [5]	Any single screenshot or single-run "rank" is unfalsifiable
Tools disagree with each other [6]	"Our dashboard says so" is not, by itself, proof

A market where claims are hard to falsify is a market that attracts inflated claims. This is the dynamic Fishkin warned about when he wrote that AI visibility "experts" could weaponize inconsistency the way "sketchy SEO salespeople did in decades past" [5]. It is also why Google's own guidance keeps repeating that optimizing for generative AI features is grounded in the same Search index and quality systems as ordinary SEO [1][2] – a fact that quietly deflates the mystique some vendors wrap around "proprietary GEO techniques."

The good news: honest GEO reporting is entirely possible right now. It just looks different from what most clients are shown – less like a rank report, more like a methodology disclosure. The next chapter catalogs what the fake version looks like; Chapter 3 shows the real one.

◆ SitesBrand position

We sell AI search optimization [11]. It is precisely because we do that we want the verification bar raised: agencies doing real work benefit when clients can tell the difference. Everything we ask you to demand from any agency in this book, you are welcome to demand from us first.

Chapter 2: The Vanity-Metric Checklist — Report Lines That Sound Good but Prove Nothing

Key Takeaways

- Vanity metrics in GEO reports share one property: **they cannot be falsified by you**, the client.
- The eight line items below are the most common ways an impressive-looking report can carry zero verifiable information.
- For each one, this chapter explains *why* it proves nothing (with the evidence from Chapter 1) and what to ask for instead.

A note on framing: this chapter describes report patterns and their logical weaknesses. It does not claim any named company uses them deceptively. The reasoning is ours; the underlying measurement facts carry citations.

The Eight Classic Vanity Lines

① "Your brand ranked #2 in ChatGPT for [query]" — single-run AI rankings

Why it proves nothing: AI answers are non-deterministic. SparkToro's 2,961-prompt study found less than a 1-in-100 chance of two runs producing the same brand list, and ~1-in-1,000 of the same order [5]. A one-off "rank" is a snapshot of randomness.

Ask instead: "Across how many runs, on which dates, and what percentage of runs included us?"

② The screenshot of a flattering AI answer

Why it proves nothing: given run-to-run variance [5], a determined reporter can re-run a prompt until the desired answer appears, screenshot it, and discard the misses. A screenshot is anecdote, not measurement.

Ask instead: the full run log — every run, including the ones where you *didn't* appear.

③ A proprietary "AI Visibility Score: 78/100" with no methodology

Why it proves nothing: with no engine ground truth to calibrate against [7] and commercial tools returning different answers to identical prompts [6], any composite score is only as meaningful as its disclosed methodology. An undisclosed formula can be tuned to go up.

Ask instead: the formula's inputs — prompt list, engines, run counts, weighting — and whether the score is comparable month over month.

④ "AI impressions are up 340%" — growth with no baseline or denominator

Why it proves nothing: percentage growth from an undisclosed base is arithmetic theater; tripling from 10 to 34 impressions is "+240%." Google's new Generative AI reports show impressions but currently **no clicks** [3][4], so an impressions-only line can't demonstrate business impact by itself either.

Ask instead: absolute numbers, the time window, the source (GSC export? tool? which?), and the same figure for the three prior periods.

⑤ "Mentioned across 12 AI platforms" — platform-count padding

Why it proves nothing: counting platforms says nothing about *frequency* of appearance on the platforms that matter to your buyers. One lucky run on twelve engines is twelve "mentions."

Ask instead: visibility percentage per engine for your priority prompts — the metric the research supports [5].

⑥ "We optimized 40 pages for GEO" — activity metrics dressed as outcomes

Why it proves nothing: this is an input, not a result. It may be perfectly good work — but a report that stops at activity is asking you to grade effort, not impact. (This is our reasoning; no citation needed — it's definitional.)

Ask instead: the before/after visibility and business metrics those 40 pages were supposed to move, per the agency's own plan.

⑦ "Proprietary llms.txt and AI-meta deployment" — technical mystique

Why it proves nothing: Google's official guidance explicitly states you do **not** need llms.txt, AI-specific meta tags, or Markdown page variants for its generative AI features — Google may crawl such files but gives them no special treatment [1]. Selling these as premium "GEO infrastructure" for Google visibility contradicts the vendor's own documentation.

Ask instead: which engine, specifically, consumes each deliverable — with the vendor documentation that says so.

⑧ "AI traffic is converting at 9x" — attribution overreach

Why it proves nothing: as Omnia's analysis notes, AI-driven brand exposure often produces **no trackable click at all** — people see a recommendation, then arrive later by brand search or direct visit [7]. Precise-sounding AI-conversion multipliers built on tiny referral segments are statistically fragile, and the influence that matters most is largely invisible to click-based attribution [7].

Ask instead: the actual segment definition, session counts (not just rates), and an honest statement of what *can't* be attributed.


The Pattern Behind All Eight

Every line above fails the same three-part test. Tape this to your monitor:

The Falsifiability Test 1. **Could I reproduce this number myself** with the stated method? (If no method is stated — fail.) 2. **Does it have a denominator and a baseline?** (Growth from what? Out of how many runs?) 3. **Does it connect to a business quantity I already track** — leads, bookings, revenue, sales-cycle inputs? (If it lives only inside the report — fail.)

A report line that passes all three is worth paying for. A report full of lines that pass none of them is decoration — regardless of how the engagement is actually performing underneath.

How to audit any line in your agency's GEO report



◆ SitesBrand position

Our own public positioning warns clients away from guaranteed rankings, bulk backlinks, and ranking-only reports [10] — this chapter simply extends that standard to the AI era. If your current agency's report is built from lines ①–⑧, the next two chapters are your exit ramp — or theirs.

Chapter 3: What a Real GEO Report Contains

Key Takeaways

- A trustworthy GEO report is recognizable by five components: **disclosed prompt methodology, repeated-run visibility percentages, named competitors, engine-native data where it exists, and a business-outcome anchor.**
- The research supports one core visibility metric: **percentage of runs in which your brand appears**, measured across dozens-to-hundreds of repetitions [5].
- Business impact should be reported the way outcome-driven case studies are written: **absolute units + percentage change + timeframe** — illustrated in this chapter with the real published figures from SitesBrand's Trackora, Finovo, and Healthify engagements [9].

3.1 Component One: Disclosed Prompt Methodology

The foundation of every honest GEO number is a **frozen, documented prompt set**. The report should state:

- **Which prompts** are tracked (the actual text, not "50 relevant queries")
- **Why those prompts** — mapped to your buying journey, agreed with you at kickoff
- **Which engines**, which geography/account state, and **how many runs per prompt per period**
- **What changed** in the set this period (additions/removals are fine; silent swapping is not)

Why run counts matter so much: with less than a 1-in-100 chance of identical answers across two runs [5], the difference between "we ran it once" and "we ran it 60 times" is the difference between anecdote and measurement. This is also how the serious end of the tooling market works — running each prompt multiple times and averaging is exactly how tracking platforms attempt to tame the variance [5][6].

3.2 Component Two: Visibility Percentage, With Competitors

The headline metric should read like this (*illustrative format — the format is the point, not these numbers*):

"Best [category] for [use case]" — ChatGPT, 60 runs, June 2026: **You: 42% of runs · Competitor A: 55% · Competitor B: 31%** · (May: You 28%)

Four properties make this line trustworthy where Chapter 2's lines were not:

1. It has a **denominator** (60 runs — falsifiable).
2. It is a **percentage across repetitions**, the metric SparkToro's research identifies as the only reasonable one given AI inconsistency [5].
3. It includes **named competitors**, so "up" has a reference frame.

4. It shows the **prior period**, so trend is visible against the same methodology.

3.3 Component Three: Engine-Native Data Where It Exists

For Google's AI surfaces, first-party data now exists: the **Search Generative AI performance reports** in Search Console provide impressions within AI Overviews and AI Mode by page, country, device, and date [3][4]. A real report:

- includes these GSC exports directly (or grants you *read access to your own Search Console* — you should have it regardless);
- labels the known limitation honestly: **no click data yet** [4];
- never presents tool-estimated numbers as if they were engine-native ones.

If your property doesn't have the report yet (Google is rolling it out to a subset of sites first [3]), the report should say that too, rather than substituting a look-alike metric without comment.

3.4 Component Four: Honest Attribution Boundaries

Because AI-influenced buyers often arrive with no trackable AI referral — via brand search or direct visits [7] — a truthful report treats attribution as a range, not a point:

- **Directly measurable:** referral sessions from chatgpt.com, perplexity.ai, gemini.google.com, etc., with raw session counts.
- **Correlated:** branded-search volume and direct-traffic movement alongside visibility gains — presented as correlation, labeled as such.
- **Unknown:** a plain statement of what cannot be attributed. An agency that never says "we can't measure that part" is describing a measurement system that doesn't exist [7].

3.5 Component Five: The Business-Outcome Anchor

Every visibility metric should ladder up to a quantity your CFO already recognizes. The reporting format we use for this at SitesBrand — and the format we suggest you demand from anyone — is the **case-study format**: absolute operational units, plus percentage change, plus timeframe, plus the client's own words. Here is what that looks like with our real published engagement data [9]:

Trackora — reported outcomes (10-week engagement)




Figure 1: Trackora (logistics, 10-week engagement) — percentage outcomes anchored to 8,347 shipments brought into real-time view and a 96.2% on-time rate. Chart built from the figures published at sitesbrand.com/case-studies [9].

Finovo — reported outcomes (8-week engagement)




Figure 2: Finovo (fintech, 8-week engagement) — outcomes anchored to 1,246 invoices tracked in a single workflow and an 18-day average collection cycle [9].

Healthify — reported outcomes (12-week engagement)

3,892 appointments · 2,341 new-patient actions · 91% retention · Source: sitesbrand.com/case-studies




Figure 3: Healthify (healthcare, 12-week engagement) — outcomes anchored to 3,892 appointments supported, 2,341 new-patient actions, and 91% patient retention [9].

Notice what makes this format resistant to inflation — and steal these properties for your own agency's reporting template:

Operational volume behind the outcomes — absolute units, not percentages

Every % claim above is anchored to countable units. Source: sitesbrand.com/case-studies (accessed 2026-07-04)




Figure 4: The absolute units behind the percentages. A percentage without a unit count is arithmetic theater (Chapter 2, line ④); a percentage anchored to 8,347 real shipments is checkable against the client's own systems [9].

1. **Absolute units first** (8,347 shipments; 1,246 invoices; 3,892 appointments) — countable in the *client's* systems, not just the agency's dashboard [9].
2. **Percentages second**, derived from those units.
3. **A timeframe** (10, 8, 12 weeks) — so "growth" has a window.

4. **The client's voice** — each SitesBrand case study carries a client quote, and the page discloses that data is anonymized for confidentiality [9].

(Transparency note: these are engagement outcomes spanning strategy, UX, development, automation, and content work as described on the case-studies page [9] — we present them here as a reporting-format example, not as GEO-specific results or as promised outcomes for any other business.)

3.6 The Report Skeleton, Assembled

1. METHODOLOGY BOX	prompt set (verbatim) · engines · runs per prompt · account/geo state · what changed this period
2. VISIBILITY	% of runs per priority prompt, per engine, vs named competitors, vs prior period
3. ENGINE-NATIVE DATA	GSC Generative AI report exports (impressions; "no clicks yet" labeled) + classic GSC/GA4
4. ATTRIBUTION	measured referrals (raw counts) · correlations (labeled) · explicit "unknown" statement
5. BUSINESS ANCHOR	units + % + timeframe, case-study format
6. WORK LOG	what was shipped, mapped to which metric it was meant to move
7. NEXT-PERIOD BETS	what should move next, by how much, so the next report is checkable against this one

Section 7 is the quiet killer: an agency that writes down *in advance* what should move next period is an agency that expects to be measured. That single habit filters more vendors than any tool audit. (*Our reasoning — no citation.*)

◆ SitesBrand position

Our AI Search Optimization service measures "queries, pages, leads, and AI-search opportunities" and runs a four-phase loop of audit, mapping, build, and monitoring [11], inside the 90-day Discovery → Strategy → Execution cadence described in our SEO-services guide [10]. The skeleton above is the reporting standard we hold that work to — bring it to your current provider and watch how they react.

Chapter 4: Five Questions to Ask Any Agency (Before Signing — or Mid-Contract)

Key Takeaways

- Each question below has a **good-answer pattern** and a **red-flag pattern**. You don't need technical knowledge to score the answers — you need to watch for falsifiability.
 - These work equally well in a sales call for a new agency or a quarterly review with your current one.
 - Ask them in order; each builds on the last.
-

Question 1 — "Which exact prompts will you track, how many runs per prompt, and will I see the raw run log?"

Why it works: it forces the methodology conversation immediately. Single-run results are statistically meaningless given documented AI answer variance [5], so the run count *is* the credibility of every number that follows.

Good answer sounds like: "We'll agree a prompt set with you at kickoff, freeze it, run each prompt N times per engine per month, and you get the raw log — including runs where you don't appear."

Red flag sounds like: "Our platform handles that automatically." (Which platform? How many runs? Tools disagree with each other on identical prompts [6] — 'the tool says so' is the beginning of the answer, not the end.)

Question 2 — "When your tool and another tool disagree, what's your ground truth?"

Why it works: there is no engine-provided console for most AI platforms [7], and three tools given the same prompts return three different answers, per the CEO of the AI-native services agency /prompt, quoted in trade press [6]. An honest agency has thought about this; a vanity-metrics agency has never been asked.

Good answer: "There is no absolute ground truth for non-Google engines right now. We mitigate with repeated runs, a stable methodology, and by treating trends within one consistent method as the signal — and we'll show you Google's own Search Console Generative AI data for Google surfaces [3][4]."

Red flag: "Our tool is the most accurate one." (No published benchmark exists to support such a claim [6].)

Question 3 — "Which business metric will this move, by roughly when — and what will you write down *now* so we can check it later?"

Why it works: it converts marketing language into a falsifiable commitment. It's the same discipline as outcome-format case reporting — units, percentages, timeframes [9] — applied *forward*.

Good answer: a specific, dated, written expectation ("qualified demo requests from organic + AI channels, visible movement inside 90 days, baseline captured in week 1"), plus honesty that external factors add noise.

Red flag: "AI visibility is a long-term brand play — it can't really be tied to revenue." (Long-term is real; *unmeasurable forever* is a tell. Even brand-level effects should show up in branded search and direct traffic, reported as labeled correlation [7].)

Question 4 — "What can't you measure? Say it out loud."

Why it works: the honest answer is well documented — AI-influenced buyers often leave no trackable referral, arriving later via brand search or direct [7]; Google's AI reports currently lack click data [4]; non-Google engines provide no first-party reporting [7]. An agency that claims full-funnel AI attribution is describing a system that does not currently exist.

Good answer: a fluent, specific list of blind spots, unprompted — plus how they bracket the uncertainty (correlation reporting, labeled estimates).

Red flag: "Our attribution covers everything end to end."

Question 5 — "Show me a report from another client (redacted). Walk me through what each number would let *them* verify."

Why it works: it applies Chapter 2's Falsifiability Test to a real artifact in real time. You're not judging design polish — you're checking: does each line have a method, a denominator, a baseline, a business anchor?

Good answer: a report that resembles Chapter 3's skeleton — methodology box, run-based visibility percentages with competitors, engine-native data where it exists, an explicit unknowns section, forward bets.

Red flag: a wall of screenshots, a proprietary score with no formula, platform-count padding — the vanity checklist, laminated.

Scoring What You Hear

	Q1 Methodology	Q2 Ground truth	Q3 Business tie	Q4 Limits	Q5 Sample report
Pass	Frozen prompts, stated run counts, raw log offered	Admits no absolute truth; consistent method as signal	Written, dated, checkable	Fluent about blind spots	Passes Falsifiability Test
Fail	"The tool handles it"	"Ours is the accurate one"	"Can't tie to revenue"	"We measure everything"	Screenshots + mystery score

Three or more fails: walk away (or, mid-contract, put the Chapter 3 skeleton in writing as your new reporting requirement and see what happens at the next review).

(The scoring rubric is our professional judgment, built on the cited measurement facts — treat the thresholds as a starting point, not statistics.)

◆ SitesBrand position

We publish warnings against guaranteed rankings and ranking-only reports on our own site [10], we disclose our service's measurement scope publicly [11], and our case studies anchor percentages to countable units with client quotes [9]. Ask us these five questions on a call — sitesbrand.com — we built this book so that you would.

The One-Page GEO Accountability Checklist

Print this page. Bring it to every proposal review and every quarterly report.

◆ Before you sign

- ☐ Prompt set agreed **in writing**, verbatim, mapped to our buying journey
- ☐ Runs per prompt per engine stated as a number (not "regularly")
- ☐ Raw run logs promised — including runs where we don't appear
- ☐ We hold owner access to our own Search Console & GA4
- ☐ Written 90-day expectation: which business metric, moved how much
- ☐ Agency stated, unprompted, what it **cannot** measure

◆ Every report must contain

- ☐ Methodology box (prompts · engines · run counts · changes this period)
- ☐ Visibility as **% of runs**, not one-off "ranks" or screenshots
- ☐ Named competitors on the same prompts, same method
- ☐ Prior-period comparison under an unchanged methodology
- ☐ GSC Generative AI export for Google surfaces (impressions; "no clicks yet" labeled)
- ☐ Absolute units next to every percentage (the denominator rule)
- ☐ Work log mapped to the metric each item was meant to move
- ☐ Next-period bets written down — checkable at the next review

◆ Walk-away red flags

- ☐ Single screenshots of AI answers presented as results
- ☐ "AI rank #N" without run counts
- ☐ Proprietary visibility score with undisclosed formula
- ☐ Growth percentages with no baseline or absolute numbers
- ☐ "Mentioned on 12 platforms" platform-count padding
- ☐ llms.txt / "AI meta tags" sold as premium Google-GEO infrastructure
- ☐ End-to-end AI attribution claimed with confidence
- ☐ Guaranteed placements in AI answers

◆ The five questions (Ch. 4)

1. Which exact prompts, how many runs, where's the raw log?
2. When tools disagree, what's your ground truth?
3. Which business metric moves, by when — written down now?
4. What can't you measure? Say it out loud.
5. Walk me through a redacted client report, line by line.

Score: 3+ red-flag answers = walk away.

From **The GEO Accountability Checklist** © 2026 SitesBrand · Full sources & methodology: sitesbrand.com · Free audit: sitesbrand.com/free-audit

Sources

All sources were opened and read on the access date shown. Inline citation numbers throughout the book refer to this list.

[1] Google Search Central — *Google's Guide to Optimizing for Generative AI Features on Google Search* (official documentation; states that generative AI features are grounded in the core Search index via retrieval-augmented generation, and that `llms.txt`, AI-specific markup, and Markdown page variants are not needed). <https://developers.google.com/search/docs/fundamentals/ai-optimization-guide> — Google for Developers. Accessed 2026-07-04.

[2] Google Search Central Blog — *A new resource for optimizing for generative AI in Google Search* (May 2026 launch announcement of [1]). <https://developers.google.com/search/blog/2026/05/a-new-resource-for-optimizing> — Google for Developers. Accessed 2026-07-04.

[3] Google Search Central Blog — *Introducing Search Generative AI performance reports in Search Console* (June 3, 2026; impressions for AI Overviews/AI Mode by page, country, device, date; rollout to a subset of websites first). <https://developers.google.com/search/blog/2026/06/gen-ai-performance-reports> — Google for Developers. Accessed 2026-07-04.

[4] Google Search Console Help — *Generative AI performance report (Search)* (report documentation; click data not included in the current version). <https://support.google.com/webmasters/answer/16984139> — Google Support. Accessed 2026-07-04.

[5] Rand Fishkin (with Patrick O'Donnell) — *NEW Research: AIs are highly inconsistent when recommending brands or products; marketers should take care when tracking AI visibility* (January 28, 2026; 600 volunteers, 2,961 prompts across ChatGPT, Claude, and Google AI; <1/100 chance of identical brand lists across runs; ~1/1,000 for identical order; recommends visibility % across dozens-to-hundreds of runs; "weaponize" quote). <https://sparktoro.com/blog/new-research-ais-are-highly-inconsistent-when-recommending-brands-or-products-marketers-should-take-care-when-tracking-ai-visibility/> — SparkToro. Accessed 2026-07-04.

[6] Digiday — *Marketers question expensive AI visibility tools as inconsistent results fuel skepticism* (May 1, 2026; Paul Dyer "three different tools... three different answers" quote; Ryan Mason "just a benchmarker" quote; lack of standardized benchmarks; Profound pricing \$99–\$399/mo; Ahrefs Brand Radar \$129–\$449/mo). <https://digiday.com/marketing/marketers-question-expensive-ai-visibility-tools-as-inconsistent-results-fuel-skepticism/> — Digiday. Accessed 2026-07-04.

[7] Omnia — *Generative Engine Optimization Challenges: What Makes GEO Hard, What the Data Reveals, & Best Practices* (May 26, 2026; "GEO produces no equivalent of Google Search Console"; AI-driven brand appearances rarely produce trackable clicks; attribution built on click events misses the shaping of consideration sets). <https://www.useomnia.com/blog/generative-engine-optimization-challenges> — Omnia. Accessed 2026-07-04.

[8] Aggarwal, P., et al. — *GEO: Generative Engine Optimization* (the academic paper that introduced the GEO term and benchmark). <https://arxiv.org/abs/2311.09735> — arXiv. Accessed 2026-07-04.

[9] SitesBrand — *Case Studies* (Trackora: 8,347 shipments in real-time view, +142% traffic, +28% on-time delivery, -37% delivery exceptions, 96.2% on-time rate, 10 weeks; Finovo: 1,246 invoices in one workflow,

+45% revenue, +32% paid invoices, -26% collection time, 18-day average collection, 8 weeks; Healthify: 3,892 appointments, 2,341 new-patient actions, +67% new patients, +38% appointments, +54% revenue, 91% retention, 12 weeks; client quotes; anonymization disclosure). <https://www.sitesbrand.com/case-studies> — SitesBrand. Accessed 2026-07-04.

[10] SitesBrand — *What to Expect From SEO Services in 2026* (90-day Discovery/Strategy/Execution cadence; warnings against guaranteed rankings, bulk backlinks, and ranking-only reports; "business outcomes, not vanity metrics" positioning). <https://www.sitesbrand.com/resources/blog/what-to-expect-from-seo-services-in-2026> — SitesBrand. Accessed 2026-07-04.

[11] SitesBrand — *AI Search Optimization* service page (deliverables: AEO/GEO strategy, entity and schema mapping, answer-ready content, AI visibility roadmap; four-phase process; measurement of queries, pages, leads, and AI-search opportunities). <https://www.sitesbrand.com/services/ai-search-optimization> — SitesBrand. Accessed 2026-07-04.

[12] Insightland — *AI Search Optimization (GEO)* service page (example of an existing how-to GEO framework: the CAT — Content, Authority, Technology — playbook). <https://insightland.org/services/ai-search-optimization-geo/> — Insightland. Accessed 2026-07-04.

[13] OptimizeGEO — *Step-by-Step Guide to GEO in 2026* (example of an existing how-to GEO framework and platform: diagnostics, insights, recommendations). <https://www.optimizegeo.ai/blog/step-by-step-guide-to-geo-2026> — OptimizeGEO. Accessed 2026-07-04.

[14] SitesBrand — *Resources: Guides* (listing of the 58-page *Ultimate 2026 Guide to Modern Search Optimization*, which covers SEO, AEO, GEO, AIO, and SXO as one system — the how-to companion to this book). <https://www.sitesbrand.com/resources/guides> — SitesBrand. Accessed 2026-07-04.

Verification Notes

- Statements labeled "our reasoning," "our position," or "our professional judgment" in the text are the authors' analysis and are deliberately **not** attached to citations.
- Three claims encountered during research were **excluded** because we could not verify them at their original source: an industry-wide AI-visibility spending estimate, a third-party brand-visibility persistence statistic, and a week-over-week citation-volatility figure. Where a claim could not be verified, it does not appear in this book.
- The Trackora, Finovo, and Healthify outcomes span full engagements (strategy, UX, development, automation, content) as described at [9]; they are used here as a reporting-format example, not as GEO-specific results or promised outcomes.

****The next move is yours.**** Take the one-page checklist into your next agency conversation — including one with us. ****Free audit:**** [sitesbrand.com/free-audit](https://www.sitesbrand.com/free-audit) · ****Strategy call:**** [sitesbrand.com](https://www.sitesbrand.com)



**If you can't falsify it,
don't pay for it.**

GEO is real. GEO accountability is optional — unless you, the client, demand it. This book gave you the vanity-metric checklist, the real report skeleton, the five questions, and the one-page checklist to enforce all of it.



Hold every agency to it.

Including us → sitesbrand.com/free-audit

PSYCHOLOGY-DRIVEN AI SEARCH & GROWTH AGENCY
FOUNDED BY HASSAM SHABBIR · [SITESBRAND.COM](https://sitesbrand.com)